

A close-up photograph of a black ink bottle with a gold cap, tilted at an angle. It is pouring a dark, viscous liquid, likely black ink, into a body of water. The water surface is dark and shows concentric ripples emanating from the point of impact. The background is blurred, focusing attention on the bottle and the liquid.

# Functions, Integrity

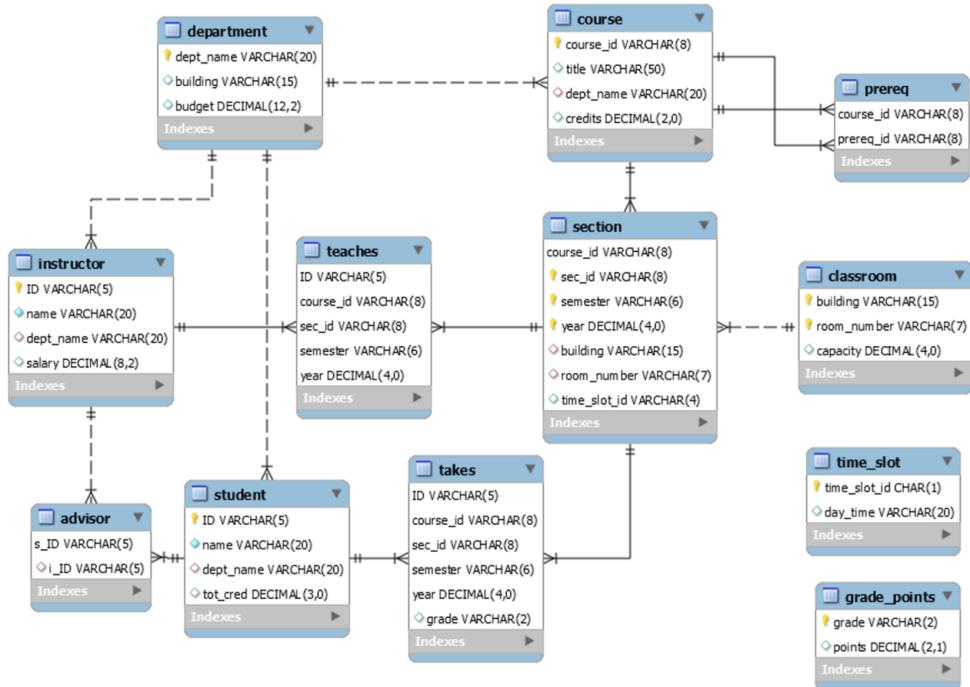
CST 363

# Topics

- `SELECT DISTINCT`
- *Scalar* functions
- `LIKE` , `BETWEEN` , `TUPLE` predicates
- `NULL` values
- Constraints
  - `CHECK` constraints
  - `UNIQUE`
  - Foreign keys
- Cascading actions

Brief Entity Relationship (ER) Diagram Aside

# Example ER Diagram for University Database



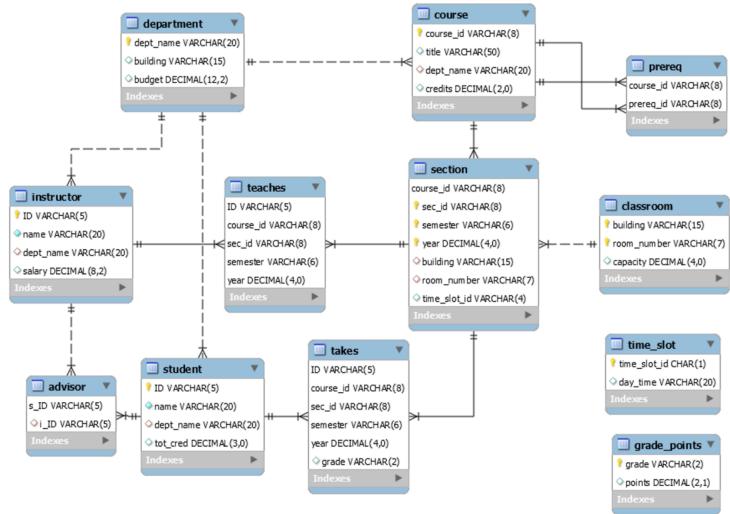
# ER Diagram – Crow's Foot Notation

+	one
←	many
++	one (and only one)
O+	zero or one
↖	one or many
○↖	zero or many

# ER Diagram – Weak vs Strong Relationships

- Dotted line – **weak / non-identifying** relationship
  - The *child* table references the *parent's* primary key as a foreign key, but the child's own primary key does not depend on it. The child still has its own identity.
  - FK exists, but not part of PK
- Solid line – strong / identifying relationship
  - FK is part of PK

# Example ER Diagram for University Database (Again)



## Identifying (solid)

### course → section

- A section is identified by (course\_id, sec\_id, semester, year) (parent key participates).

- In your diagram, section includes course\_id as part of its **PK** → identifying.

## Non-identifying (dotted)

### department → course

- course has its own **PK** (course\_id).
- dept\_name is **FK** but not part of course **PK** → non-identifying.

# SELECT DISTINCT

- SQL allows duplicates in relations as well as in query results.
- To force the elimination of duplicates, insert the keyword `DISTINCT` after `SELECT` .
  - `DISTINCT` applies to the entire selected row (all selected columns), not each column independently.
- Find the department names of all instructors, and remove duplicates

```
SELECT dept_name  
FROM instructor;
```

instructor
dept_name
Comp. Sci.
Finance
Physics
History
Physics
Comp. Sci.
History
Finance
Biology
Comp. Sci.
Elec. Eng.

```
SELECT DISTINCT dept_name  
FROM instructor;
```

instructor
dept_name
Comp. Sci.
Finance
Physics
History
Biology
Elec. Eng.

# "Scalar" Functions

A scalar function returns one value per row (or per call) – not a table/result set.

- SQL supports a variety of string operations such as
  - concatenation
    - `CONCAT(last_name, ', ', first_name)`
  - converting from upper to lower case (and vice versa)
    - `UPPER(name), LOWER(name)`
  - finding string length, extracting substrings, find a substring etc.
    - `LENGTH(name), SUBSTRING(name, 3, 4)`
    - `SELECT POSITION('SU' IN 'CSUMB');` → returns 2
    - SQL starts indexing with 1, not 0

## Scalar Functions (cont.)

- Numeric functions include
  - `ROUND(23.66666, 1)` → 23.7
  - `ROUND(23.66666)` → 24

# LIKE Predicate

- The predicate `LIKE` uses patterns to match strings that are described using two special characters:
  - percent ( `%` ). The `%` character matches any substring.
  - underscore ( `_` ). The `_` character matches any character.
- Find the names of all instructors whose name includes the substring "dar".

```
SELECT instructor_name  
FROM instructor  
WHERE instructor_name LIKE '%dar%';
```

# LIKE Predicate - Examples

- `LIKE '100\%' ESCAPE '\'` matches the string 100%
- `LIKE 'Intro%'` matches any string beginning with "Intro".
- `LIKE '%Comp%'` matches any string containing "Comp" as a substring.
- `LIKE '___'` matches any string of exactly three characters.
- `LIKE '___%'` matches any string of at least three characters

# BETWEEN Predicate

- SQL includes a `BETWEEN` comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000

```
SELECT instructor_name  
FROM instructor  
WHERE salary BETWEEN 90000 AND 100000;
```

equivalent to:

```
SELECT instructor_name  
FROM instructor  
WHERE salary >= 90000 AND salary <= 100000;
```

# Tuple comparison

```
SELECT instructor_name, dept_name  
FROM instructor  
WHERE (ID, dept_name) = ('45565', 'Comp. Sci.');
```

equivalent to:

```
SELECT instructor_name, dept_name  
FROM instructor  
WHERE ID = '45565' AND dept_name = 'Comp. Sci.';
```

# Null Values

- The predicate `IS NULL` can be used to check for null values.
- Example: Find all instructors whose salary is `NULL`.

```
SELECT instructor_name  
FROM instructor  
WHERE salary IS NULL;
```

- The predicate `IS NOT NULL` succeeds if the value on which it is applied is not null.
- The result of an expression involving `NULL` is `NULL`
- *Example:* `5 + NULL` returns `NULL`

## Null Values (Cont.)

- SQL treats as unknown the result of any comparison involving a null value (other than predicates `IS NULL` and `IS NOT NULL` ).
  - Example: `5 < NULL` or `NULL < NULL` or `NULL = NULL`
- The predicate in a `WHERE` clause can involve Boolean operations (`AND`, `OR`); thus the definitions of the Boolean operations need to be extended to deal with the value **unknown**.
  - `AND` : (true AND unknown) = unknown,  
(false AND unknown) = false, (unknown AND unknown) = unknown
  - `OR` : (unknown OR true) = true, (unknown OR false) = unknown, (unknown OR unknown) = unknown
- Result of `WHERE` clause predicate is treated as false if it evaluates to unknown
- **Note:** `= NULL` never works; use `IS NULL`.

# Integrity Constraints

- An integrity constraint is a condition on data.
- If an integrity constraint doesn't hold, there is a problem (or inconsistency) with the data.
- Integrity constraint can involve
  - invalid key (primary key or foreign key) value
  - invalid data value (negative salary value for example)
  - uniqueness (2 different students same social security number)

You define the integrity constraints; the database enforces them.

# Recall: Constraints

- PRIMARY KEY = UNIQUE + NOT NULL
- FOREIGN KEY = must reference an existing parent key
- CHECK = row-level predicate must be true

## SQL CHECK constraints (Example 1 of 3)

```
CREATE TABLE department (
    dept_name VARCHAR(20),
    building VARCHAR(15),
    budget NUMERIC(12, 2) CHECK (budget > 0),
    PRIMARY KEY (dept_name)
);
```

- A department's budget must be greater than 0

## SQL CHECK constraints (Example 2 of 3)

```
CREATE TABLE section (
    course_id VARCHAR(8),
    sec_id VARCHAR(8),
    semester VARCHAR(6) CHECK (
        semester IN ('Fall', 'Winter', 'Spring', 'Summer')
    ),
    ...
);
```

- Semester must be one of Fall/Winter/Spring/Summer.

## SQL CHECK constraints (Example 3 of 3)

```
CREATE TABLE course (
    course_id VARCHAR(8),
    title VARCHAR(50),
    dept_name VARCHAR(20),
    credits NUMERIC(2, 0),
    CHECK (
        dept_name ◇ 'Comp. Sci.' OR
        course_id LIKE 'CS-%'
    )
);
```

- If dept\_name is "Comp. Sci." then course\_id must start with "CS-"
- ◇ can also be written in PostgreSQL as ≠

# SQL UNIQUE constraints

```
CREATE TABLE department (
    dept_name VARCHAR(20) PRIMARY KEY,
    dept_abrv VARCHAR(4),
    building VARCHAR(15),
    budget NUMERIC(12, 2),
    UNIQUE (dept_abrv)
);
```

No two rows can have the same non-null `dept_abrv` values

Things to remember:

- `UNIQUE` constraints look like primary key constraints except unique columns can be a null value, primary key's do not allow null values.
- A table can have many `UNIQUE` constraints
- `NULL = x` is always false, so `UNIQUE` only cares about non-null values

# SQL FOREIGN KEY constraints

A foreign key column(s), if not null, must match a candidate key (typically the PK) of the referenced table.

```
CREATE TABLE course (
    ... ,
    dept_name VARCHAR(20),
    FOREIGN KEY (dept_name) REFERENCES department (dept_name),
    ...
);
```

- `INSERT INTO course VALUES ( ... , NULL, ... );` → OK
- `INSERT INTO course VALUES ( ... , 'Comp. Sci.', ... );` → OK
- `INSERT INTO course VALUES ( ... , 'CompSci.', ... );` → Constraint violation, misspelled

# Cascading Actions in Referential Integrity

- When a referential-integrity constraint is violated, the normal procedure is to reject the action that caused the violation.
- An alternative, in case of `DELETE` or `UPDATE` is to `CASCADE`.
  - If the department is deleted, delete all courses related to that department.
  - If the department name is updated in the department table, change the foreign key values in the course's table.

```
CREATE TABLE course (
    ...,
    dept_name VARCHAR(20),
    FOREIGN KEY (dept_name) REFERENCES department(dept_name)
        ON DELETE CASCADE
        ON UPDATE CASCADE,
    ...
);
```

## Cascading Actions in Referential Integrity (cont.)

- Instead of `CASCADE` , we can use `SET NULL` or `SET DEFAULT` :
  - `SET NULL` : If the referenced row is deleted or updated, the foreign key is set to `NULL` .
  - `SET DEFAULT` : If the referenced row is deleted or updated, the foreign key is set to its default value (if a default is defined).